

DRILLER at GenSIE 2026

Ariel González Gómez¹, Daniel Valdés Pérez¹

¹Havana University, Havana, Cuba

Abstract

This paper describes DRILLER, the system submitted to GenSIE 2026 for Spanish schema-guided information extraction. GenSIE requires systems to extract grounded JSON objects from Spanish texts under variable JSON Schemas provided at inference time. In this setting, the main difficulty is not only producing syntactically valid JSON, but also interpreting each field according to a task-specific schema whose names, descriptions, types, and constraints may change from one instance to the next. DRILLER addresses this problem through three submitted pipelines built around a common schema-guided extraction architecture. The system renders schemas in a model-readable form, constrains generation with a strict structured response format, and enriches prompts with field-aware retrieval over complementary demonstrations. One variant introduces temporary top-level reasoning/value wrappers so that the model can deliberate locally about evidence before producing each final field value. A third variant applies heterogeneous self-consistency by combining direct and reasoning-augmented extractors, then selecting a final prediction through schema-aware aggregation. Across the submitted pipelines, DRILLER emphasizes schema interpretation, grounded extraction, conservative postprocessing, and compatibility with the model-agnostic hosted inference protocol of the shared task.

Keywords

Spanish information extraction, schema-guided extraction, retrieval-augmented generation, structured generation, self-consistency

1. Introduction

Schema-guided information extraction requires systems to adapt to an output contract that may be unknown before inference time. Instead of extracting a fixed inventory of slots, the model receives a task-specific schema and must decide what each field means for the current source text. This setting is increasingly common when language models are used behind APIs, form-filling workflows, or data integration tools: the schema is not just a validator, but part of the instruction.

GenSIE 2026 formalizes this problem for Spanish information extraction [1]. Each instance provides a Spanish source text, an extraction instruction, and a target JSON Schema; the system must return a grounded JSON object that conforms to that schema. The task is part of IberLEF 2026 [2] and uses a hosted, model-agnostic evaluation protocol, so a submitted system must respect the model selected by the evaluator rather than relying on a privately chosen backend. This makes inference-time methods such as schema presentation, retrieval, constrained generation, and aggregation central to the submission.

The main difficulty is broader than JSON validity. A model can return a syntactically valid object and still fail because it has misread the schema. A string field may require a copied mention, a normalized value, or a short grounded summary depending on its description. An enum may require mapping textual evidence to labels that do not appear literally. A nullable field may invite a plausible answer from world knowledge even when the passage does not support it. Arrays and nested objects add further ambiguity about completeness and item identity. In these cases, the shape of the JSON is necessary but not sufficient; the model must align each value with the current field semantics.

This field-interpretation problem is amplified by variable schemas. If the schema were fixed, field semantics could be learned from training data or encoded in task-specific rules. In GenSIE, field names, descriptions, type constraints, enum labels, nullability, and nesting are all runtime evidence. The extractor must use them to decide which text spans are relevant, when normalization is expected, when absence should be represented as JSON `null`, and how much structure to return.

Our team submitted three pipelines to the challenge. We refer to enriched-schema-rag as Schema-RAG, the direct retrieval-augmented extractor. We refer to enriched-inline-reasoning-rag as Reasoned-RAG, which uses the same retrieval path but temporarily wraps top-level fields as reasoning/value objects during generation. We refer to mixed-extractors-self-consistency-rag as Mixed-SC, which combines direct and reasoning-augmented trials through heterogeneous self-consistency and schema-aware aggregation. These aliases are used throughout the paper.

All three pipelines share a schema-guided extraction spine. DRILLER renders the target schema in a compact Pydantic-style view for the prompt, sends a strict JSON Schema response format to constrain generation, retrieves up to two field-aware demonstrations from a curated few-shot corpus, and applies deterministic postprocessing before returning the final object. The pipelines differ in whether they use a temporary reasoning scaffold and whether they aggregate multiple candidates. The method sections describe these mechanisms in detail.

Schema-readable prompting separates two uses of the schema. The prompt-side view is a communication layer that foregrounds field names, descriptions, types, enums, nullability, and nesting. The generation schema is an interface constraint that encourages the model to return a complete JSON object. This separation is useful because language models may understand typed, class-like representations more readily than raw JSON Schema, while the evaluator still expects the original JSON shape.

Field-aware few-shot retrieval addresses the fact that examples can help or hurt depending on how they relate to the current schema. DRILLER first looks for demonstrations with the same normalized schema. When exact structural matches are unavailable, it falls back to field-level retrieval with type-compatible matching and a complementary pair objective. The goal is to show relevant extraction behavior, including contrasting outcomes such as present versus absent values, without treating example content as evidence for the current instance.

Reasoned-RAG adds an internal reasoning interface for fields whose values require evidence assessment, null decisions, or schema-specific label mapping. During generation, each top-level field is represented as an object with a reasoning string and a value of the original field type. After generation, the reasoning strings are discarded and only the values are returned. Mixed-SC then uses both direct and reasoned extractors as heterogeneous candidate generators, aggregating their final-shape JSON outputs with recursive schema-aware rules for scalars, nulls, objects, and arrays.

The contributions of this work are:

- a modular schema-guided extraction architecture for variable JSON Schemas in Spanish information extraction;
- schema-readable prompting paired with strict structured generation, separating schema interpretation from output-interface enforcement;
- field-aware few-shot retrieval with same-schema matching, type-compatible field fallback, and complementary example selection;
- a temporary top-level reasoning/value transformation that supports field-local evidence assessment while preserving the final schema interface;
- heterogeneous self-consistency over direct and reasoning-augmented RAG extractors with conservative schema-aware postprocessing and aggregation.

2. Task Setting

GenSIE is a Spanish schema-guided information extraction task [1]. Each instance provides a source text, a natural-language extraction instruction, and a target JSON Schema. The source text is the only evidence for the answer. The instruction states the extraction goal, and the schema defines the output structure through field names, descriptions, primitive types, objects, arrays, enum labels, nullability, and required properties.

The required output is a JSON object grounded in the source text and conforming to the target schema. Conformance includes producing the expected object shape, using valid primitive and enum

Table 1
Submitted DRILLER pipelines.

Alias	Schema view	Reasoning wrapper	Requested trials	Final output
Schema-RAG	Pydantic	No	1	Direct
Reasoned-RAG	Reasoned Pydantic	Top-level	1	Unwrap + direct
Mixed-SC	Direct and reasoned	Reasoned trials	4	Schema-aware SC

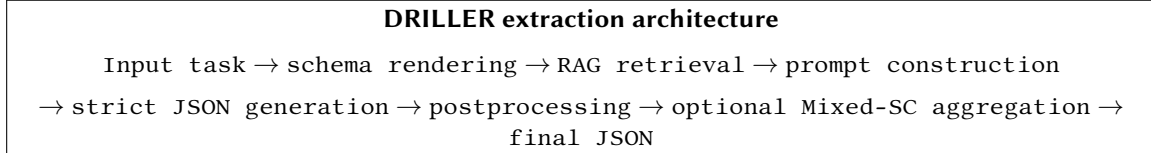


Figure 1: End-to-end architecture shared by the submitted pipelines. Aggregation is active only for Mixed-SC; the two single-pass pipelines return the postprocessed extraction directly.

values, and representing nested structures and arrays according to the schema. Grounding means that values should be supported by the current passage rather than external knowledge, retrieved examples, or plausible background assumptions.

Schemas vary across instances, so the extractor cannot rely on a fixed ontology or stable slot inventory. The same textual mention may be relevant for one schema and irrelevant for another; the same JSON type may correspond to different behaviors depending on the field description. The schema must therefore be interpreted dynamically at inference time.

Missing information must be handled conservatively. When the requested value is not supported by the source text and the schema permits null, the system should return JSON null. If evidence is present, returning null is also an error. DRILLER therefore treats nullability as a schema-guided evidence decision rather than a generic fallback.

Official dataset construction, scoring, ranking, and evaluation protocol details are specified in the GenSIE overview paper. This paper focuses on the submitted DRILLER pipelines and their inference-time mechanisms for producing grounded schema-conformant JSON under that protocol.

3. System Overview

DRILLER proposes three submitted pipelines: Schema-RAG, Reasoned-RAG, and Mixed-SC. They share a common schema-guided extraction architecture and differ only in the use of reasoning scaffolds, trial count, and final decision rule. The submitted interface names are enriched-schema-rag, enriched-inline-reasoning-rag, and mixed-extractors-self-consistency-rag, respectively. Non-registered experimental variants are outside the submitted system.

Table 1 summarizes the final pipelines using the paper aliases as primary labels. All three use retrieval-augmented prompting and strict structured JSON generation; the later method sections describe retrieval, schema rendering, reasoning wrappers, postprocessing, and aggregation in detail.

The shared architecture is a sequence of transformations from a GenSIE instance to a final JSON object. The input supplies the Spanish source text, extraction instruction, and target JSON Schema. DRILLER renders the schema for prompting, retrieves local demonstrations, constructs a Spanish extraction prompt, requests a structured JSON object from the hosted model, parses and normalizes the output, and optionally aggregates multiple candidates. The returned object always has the original challenge schema shape.

The prompt is organized as a single extraction request. It presents the instance instruction, grounding rules, retrieved examples when available, a Pydantic version of the target schema, and the source text. The rules emphasize that the source text is the only evidence, examples are demonstrations rather than facts for the new instance, enum labels must be used exactly, and unsupported nullable fields should be

returned as JSON `null`. This organization keeps the schema and passage close to generation while making the evidence boundary explicit.

Schema rendering is shared across pipelines, with a direct and a reasoned variant. In direct mode, the JSON Schema is shown as a compact Pydantic-style model: fields become typed attributes, descriptions remain near the fields they define, arrays and nested objects are readable as typed structures, and enums are displayed as literal alternatives. This prompt view is not the formal output contract. The model request also includes a strict JSON Schema response format derived from the target schema, and the final answer is checked against the evaluator-facing schema shape after postprocessing.

Retrieval is also shared. Each extraction can include up to two local few-shot demonstrations selected for schema relevance. Same-schema retrieval is used when examples have the same normalized structural schema as the current task; otherwise, DRILLER falls back to field-level matching over field names, descriptions, type classes, and complementary coverage. The selected examples are rendered in the same output style as the current extractor: direct examples for Schema-RAG, reasoning/value examples for Reasoned-RAG, and trial-specific rendering inside Mixed-SC.

Postprocessing is deterministic and conservative. It parses JSON, cleans common Unicode artifacts, normalizes strings to *Normalization Form C* (NFC), removes temporary reasoning wrappers when present, and converts the literal string `"null"` to JSON `null` only at schema positions where null is allowed. It does not infer missing facts, repair unsupported values, or reinterpret the source text after generation.

Schema-RAG is the direct single-pass pipeline. It uses the Pydantic-style prompt schema, retrieves demonstrations, requests a strict JSON object, and returns the parsed and normalized object. Because no reasoning wrapper is used, the generated top-level fields already match the final output interface.

Reasoned-RAG follows the same path but temporarily transforms each top-level field into an object with reasoning and value slots during generation. The prompt schema exposes this as a reasoned value, and the response schema requires the wrapper structure. After generation, DRILLER discards the reasoning strings and returns only the values in the original schema shape.

Mixed-SC moves from a single extraction to a heterogeneous candidate set. It requests two direct Schema-RAG-style trials and two reasoning-augmented Reasoned-RAG-style trials. Each trial performs its own retrieval, prompt construction, strict generation, and postprocessing. Aggregation receives final-shape JSON candidates, not raw model messages or reasoning traces, and combines them recursively using schema-aware rules for scalars, nulls, objects, and arrays.

Thus the submitted system keeps one common extraction spine while varying the amount of inference-time intervention. Schema-RAG returns a direct postprocessed extraction, Reasoned-RAG adds and removes a temporary field-level reasoning interface, and Mixed-SC aggregates heterogeneous candidates. The technical details of each component are given in the following sections.

4. Schema-Guided Prompting

DRILLER's prompting strategy treats the schema as part of the task semantics, not only as a formatting constraint. The prompt must help the model decide what each field asks for, while the model request must still constrain the response to a JSON object. This section describes the prompt objective, schema rendering, structured response format, and interaction with retrieval and reasoning wrappers.

The prompt objective is to make the model behave as a precise Spanish information extractor. The model is instructed to use the current source text as the only evidence, follow the extraction instruction, and interpret field meanings through the target schema. Retrieved examples are explicitly presented as demonstrations of extraction behavior, not as evidence for the current input. The answer must be only the requested JSON object, without markdown fences, explanations, or extra text.

The standard prompt is organized around the extraction workflow. It first sets the extraction role. It then gives Spanish extraction rules covering grounding, full schema coverage, exact enum labels, conservative null handling, and value normalization when requested by the field description. Retrieved examples are inserted next. Finally, the current schema is rendered in a model-readable form followed

Table 2

Schema roles in DRILLER. The prompt schema supports model comprehension; the generation schema constrains decoding; the final challenge schema defines the returned interface.

Schema role	Where it is used	Function in the system
Prompt schema	Text shown inside the extraction prompt	Pydantic-style or reasoned Pydantic-style rendering that helps the model interpret field names, descriptions, types, enums, nullability, and nesting.
Generation schema	Structured response format sent with the model request	Strict JSON Schema used to constrain the generated object; in reasoning mode, temporary wrappers are inserted before generation.
Final challenge schema	Shape expected by the evaluator	Original task schema after postprocessing and optional unwrapping; reasoning scaffolds and prompt-only conventions are not returned.

by the source text. This order keeps task intent and extraction constraints visible while placing the current schema and evidence near generation.

Raw JSON Schema is well suited for validation, but it can be visually heavy inside a prompt. Nested property dictionaries, separate required lists, repeated keywords, and nullable type arrays can obscure the field descriptions that carry extraction semantics. DRILLER therefore renders the schema in a Pydantic-style view. The view is not an executable program and does not replace the target JSON Schema; it is a compact representation designed to make the current extraction contract easier for the model to read.

In this view, object schemas become class-like structures and properties appear as typed attributes. Primitive fields are rendered as familiar scalar types, arrays as list types, enums as literal alternatives, and nested objects as nested or referenced model types. Nullability is shown near the field rather than separated from the definition. Field descriptions are preserved because they often determine the correct behavior: the same type annotation can require copying, normalization, classification, or a short grounded summary depending on the description.

DRILLER distinguishes three schema roles, summarized in Table 2. The same target schema is transformed differently depending on whether it is being shown to the model, used to constrain decoding, or returned to the evaluator.

The prompt schema is a readability layer. It can use class-like declarations, readable nullable notation, and other prompt-only conventions because its purpose is communication. The generation schema has a different role: DRILLER attaches a strict JSON Schema response format to the model request so decoding is constrained toward a JSON object compatible with the requested structure.

For the submitted non-baseline pipelines, declared object properties are required during generation. This reduces field omissions, a common failure mode when a model understands the task but leaves out keys. Requiring keys does not require every value to be substantive. If the target schema permits null, the generated object may include the key with JSON `null`. The final challenge schema remains the original schema shape; generation-time requiredness is an inference constraint, not permission to fabricate unsupported content.

Reasoning wrappers create a second schema variant. In direct mode, used by Schema-RAG, the prompt schema is ordinary Pydantic-style rendering and the model produces final values directly. In reasoning mode, used by Reasoned-RAG and the reasoning-augmented trials of Mixed-SC, each top-level field is shown as a reasoned value. The temporary generation schema requires an object with reasoning and value fields for each top-level slot. The value field retains the original type, including enums, arrays, objects, and nullable alternatives.

This transformation is a schema-guided interface change rather than a request for free-form explanation. The model has a structured place to state local evidence or absence before committing to a field value, but the reasoning text is removed during postprocessing. Direct and reasoned prompting

therefore share the same final objective while exposing different intermediate interfaces.

Prompting also interacts with retrieval. In the general RAG layout, each example carries enough schema context for the model to understand why its output follows from its own source text and schema. In the same-schema layout, selected examples share the normalized target schema, so DRILLER can render the shared schema once and show compact example inputs and outputs. This saves prompt size and lets the examples emphasize contrasting field outcomes under the same structure.

Examples can show how fields are populated, when null or empty arrays are appropriate, and how enum labels or nested structures behave, but their values are not evidence for the current instance. In reasoning mode, retrieved examples are rendered with the same temporary `{reasoning, value}` structure expected from the current extractor; in direct mode, examples use final field values. Mixed-SC applies the appropriate layout separately inside each trial.

Overall, schema-guided prompting in DRILLER combines readable schema presentation, explicit grounded-extraction rules, strict structured generation, and retrieval-aware layout. The purpose is to make schema interpretation explicit while returning a complete JSON object compatible with the evaluator-facing schema.

5. Field-Aware Few-Shot Retrieval

Retrieval in DRILLER is designed around schema interpretation. In a variable-schema task, semantic similarity between source texts is not enough: two passages may be topically close but ask for different fields, while two different domains may instantiate similar extraction behavior. The submitted pipelines therefore retrieve demonstrations with attention to schema structure, extraction instruction, field descriptions, type compatibility, and complementary output behavior. The same retrieval path is used by Schema-RAG, Reasoned-RAG, and every trial of Mixed-SC.

5.1. Curated Demonstration Corpus

The retrieval pool is a local few-shot prompting (FSP) corpus of structured extraction demonstrations. Each case contains a Spanish source text, an extraction instruction, a target schema, expected values, tags describing covered extraction phenomena, and field-level rationales or descriptions that support rendering examples in either direct or reasoning-augmented form. The corpus is synthetic and curated. Development examples were inspected by task or schema family, field contrasts were identified, synthetic cases were drafted, outputs and rationales were reviewed, and the resulting resources were validated before use.

The final extraction corpus contains two cases for each of the 22 schema or task families of the Development Set, for 44 extraction cases in the local retrieval pool. The families cover a range of domains and structures, but the corpus is not intended to be a large topical nearest-neighbor database. It is a compact library of contrasted demonstrations: the useful unit is often a pair of cases showing how the same or related fields behave under different evidence conditions.

The paired organization is deliberate. A nullable field benefits from both grounded non-null values and grounded absence. An array field benefits from examples where the correct list is populated and examples where it is empty. Enum fields benefit from contrasting labels, especially when labels are inferred rather than copied literally. Numeric and date fields may require exact copying in one case and normalization in another. String fields may require direct mention extraction, near-verbatim evidence, or a short grounded summary depending on the field description.

The documented coverage includes null versus non-null values, present versus absent evidence, enum alternatives, empty versus populated arrays, arrays of simple values, arrays of objects, nested objects, direct strings, summary strings, boolean inference, numeric normalization, date normalization, bounded scores, long evidence spans, and distractor mentions. These contrasts support retrieval because few-shot examples can bias the model when they show only one outcome regime. Complementary demonstrations help teach conditional behavior: extract when evidence exists, abstain when it does not, normalize only when requested, and respect the requested type and label set.

5.2. Same-Schema Retrieval

At runtime, DRILLER first tries to retrieve examples with the same normalized schema as the current task. This stage begins by computing a normalized schema fingerprint for the target schema and for each candidate case. The normalization canonicalizes the representation while preserving the structural extraction contract: field names, object structure, array structure, primitive types, enum alternatives, required properties, and nesting are retained. Descriptions are removed or neutralized so that minor wording differences do not prevent a match. Order-insensitive elements such as `required`, `enum`, and `array-valued type` are canonicalized.

The fingerprint distinguishes schema identity from topical similarity. A medical passage and a technical passage might both contain names, dates, arrays, and classifications without sharing an extraction form. Conversely, two source texts from the same schema family may require the same field interpretation even when their values differ. Same-schema retrieval uses structural identity as a hard filter: only cases whose normalized fingerprint matches the current target schema become candidates for this stage.

Filtered candidates are then ranked semantically. DRILLER builds a global textual representation of the current task from task identity, extraction instruction, and textual representations of top-level fields. Candidate cases are represented analogously. These representations are embedded and compared by cosine similarity. The default same-schema threshold documented for the submitted system is 0.6. If at least two same-schema candidates meet the threshold, the top two are selected as same-schema demonstrations.

Same-schema examples are valuable because the model observes the exact output contract under different source texts. They can clarify what counts as evidence for a field, when null or an empty list is correct, and how enum labels or nested structures are populated. This stage also enables a compact prompt layout: the shared schema is rendered once, and examples can focus on instruction, source text, and expected output rather than repeating identical schema blocks.

5.3. Field-Level Retrieval Fallback

When same-schema retrieval does not produce enough candidates, DRILLER falls back to field-level retrieval. The target schema is decomposed into top-level field specifications. For each field, the retriever records the field name, field description, textual type representation, and a coarse type class. Candidate cases are decomposed in the same way.

Field comparisons are constrained by type compatibility. A query field is compared only to candidate fields with compatible coarse type classes. This avoids misleading alignments where descriptions share topical words but the expected output shape differs, such as matching an array of entities to a scalar summary. Nullable and non-nullable variants of the same broad type remain comparable, while an enum is not treated as compatible with an arbitrary nested object. The goal is schema analogy: find cases where a similar kind of field was interpreted and filled.

For each compatible field pair, DRILLER computes embedding similarity over field texts. Each target field receives its best compatible similarity for a given candidate case. The result is a field-similarity vector aligned with the top-level fields of the current query. Negative similarities are clipped to zero, so a poor match simply contributes no coverage. A candidate's individual relevance can be summarized by the norm of this vector, but two-example selection uses a pair-level objective.

This vector representation is more informative than a single nearest-neighbor score. A candidate may be moderately useful for all fields, or highly useful for only one part of the schema. Since the retrieval target is a pair of demonstrations, two individually imperfect candidates can be preferable when they cover different field behaviors.

If embedding retrieval fails or produces no usable candidates, the system has a lexical/schema fallback. It scores candidates using explicit features such as exact normalized schema matches, compatible field fingerprints, tag overlap, field-name overlap, shared terms, and domain-term overlap, with a small penalty for long example texts. This fallback is not the primary path, but it preserves conservative

schema-aware retrieval when embeddings are unavailable or uninformative.

5.4. Complementary Pair Selection

The submitted RAG path retrieves up to two examples. For the field-level fallback, DRILLER selects a pair using a complementarity-aware objective rather than choosing the two highest individual scores. For two candidate demonstrations c_1 and c_2 , with field-similarity vectors v_1 and v_2 , the pair score is

$$s(c_1, c_2) = \|v_1 + v_2\|_2 + \|v_1 - v_2\|_2.$$

The first term, $\|v_1 + v_2\|_2$, rewards joint relevance and field coverage. If both examples provide useful analogues for the current schema, their combined vector has large magnitude. This term aligns with the ordinary retrieval goal that demonstrations should be relevant to the task.

The second term, $\|v_1 - v_2\|_2$, rewards complementarity. If two candidates are strong on the same fields and weak on the same fields, their difference vector is small. If one is strong where the other is weak, the difference grows. This discourages near-duplicate demonstrations that reinforce one behavior while leaving other fields unexplained.

This objective matters because demonstration diversity is not merely topical diversity. Two examples from different domains can still be redundant if they both show easy scalar fields and neither illustrates nulls, arrays, or enum decisions. Conversely, two examples from a related family may be complementary if one illustrates absence and empty arrays while the other illustrates populated nested objects and label mapping. Tie-breaking favors stronger individual relevance and stable resource order, but the core decision is pair-level.

5.5. Prompt Integration

Selected examples are integrated into the extraction prompt in one of two layouts. The general RAG layout is used when examples are related but not all same-schema matches. Each example is rendered as a self-contained demonstration with its instruction, schema context, source text, and expected output. This gives the model enough context to understand how the example output follows from that example’s schema before using it analogically.

The same-schema layout is used when all selected examples share the normalized target schema. The shared schema is rendered once, alongside the role instruction, extraction rules, root description when available, and guidance that examples are demonstrations rather than evidence. The examples then omit repeated schema blocks and focus on source-to-output behavior under the common structure. This reduces prompt length and makes contrasting field outcomes more salient.

Pipeline integration depends on the current output interface. In Schema-RAG, examples are rendered in direct final-value form, matching the object shape returned after generation. In Reasoned-RAG, examples are rendered with top-level `{reasoning, value}` objects because the temporary generation interface requires that structure. The reasoning strings demonstrate a local evidence protocol, while the value slots show the values that will later be unwrapped.

Mixed-SC uses the same retrieval machinery inside each trial. Direct trials use the Schema-RAG rendering, and reasoning-augmented trials use the Reasoned-RAG rendering. Each trial performs retrieval independently. Aggregation receives only postprocessed JSON candidates in the final schema shape.

Overall, field-aware retrieval turns the local corpus into an inference-time schema interpretation aid. Same-schema retrieval exploits exact structural matches when available, field-level fallback finds type-compatible analogies, the pair objective favors complementary coverage, and prompt integration adapts examples to the direct, reasoned, and mixed pipelines.

6. Reasoning-Augmented Extraction

Reasoned-RAG adds a temporary reasoning interface to the shared extraction pipeline. Some GenSIE fields require local evidence assessment before a value can be produced: the model may need to identify relevant fragments, decide whether evidence is absent, normalize a phrase, or map cues to an enum label. Reasoning-augmented extraction gives each top-level field a structured place for that intermediate decision while preserving the final JSON interface.

In Reasoned-RAG, each top-level field of the target object is temporarily wrapped as an object with two properties:

```
{
  "reasoning": "...",
  "value": ...
}
```

The reasoning property is a string. The value property keeps the original schema of the wrapped field. If the original field is an enum, the value must be one of the enum labels; if it is an array or object, the nested structure remains there; if it is nullable, the value may be JSON `null` when the schema permits it. The wrapper changes the temporary generation shape, not the semantic target.

The submitted reasoning mode is top-level only. It wraps root object properties, not every nested property recursively. If a top-level field is itself an object or an array of objects, the nested content remains inside value according to the original field schema. This keeps the prompt and response format manageable while asking the model to deliberate at the level of the main requested slots.

The prompt-side effect is a reasoned Pydantic-style schema view. Rather than seeing a top-level field simply as a value of type T , the model sees it as a reasoned value, using the class `Reasoned[T]`. The extraction instructions ask the model to state what the field requests, identify supporting evidence or absence, and then commit to the final value. This is intended for fields involving nullability, enum classification, boolean inference, normalization, and arrays whose correct value may be empty.

The generation-side effect is equally important. The reasoning interface is part of the strict temporary response schema, so each wrapped top-level field must contain both reasoning and value. This prevents free-form reasoning from leaking outside the JSON object and keeps the final value constrained by the original field type. The wrapper adds a field-local explanation channel without relaxing the schema constraints on the answer.

After generation, postprocessing removes the wrapper. Each top-level object of the form `{reasoning, value}` is replaced by its value; reasoning strings are discarded before the output is returned or passed to self-consistency aggregation. If the expected wrapper structure is malformed and cannot be reliably unwrapped, the extraction record is treated as invalid rather than repaired semantically.

Reasoning-augmented extraction also affects RAG rendering. Examples retrieved for Reasoned-RAG are shown in the same temporary output format expected from the current extractor. A direct example with final field values would be inconsistent with a prompt that asks for top-level `{reasoning, value}` objects. Reasoning-mode examples therefore demonstrate both the evidence protocol and the placement of the final answer in the value slot, while still remaining demonstrations rather than evidence for the current instance.

7. Postprocessing and Output Normalization

DRILLER uses deterministic postprocessing to align model outputs with the final schema interface. Even with a strict JSON Schema response format, outputs can contain representation artifacts such as literal Unicode escape strings, the string `"null"` where JSON `null` was intended, or temporary reasoning wrappers from Reasoned-RAG. Postprocessing cleans these artifacts.

The first step is JSON parsing. The submitted pipelines expect a JSON object because the prompt asks for JSON only and the request includes a structured response format. If an output cannot be parsed as JSON, the trial is marked invalid rather than repaired by string heuristics. This boundary is especially important for Mixed-SC, where aggregation should combine valid candidate objects rather than fragments of malformed responses.

After parsing, DRILLER applies Unicode cleanup recursively to string values. Spanish text may include diacritics, the letter ñ, inverted punctuation marks, and other characters that can appear inconsistently across model outputs or API layers. The normalizer cleans literal escape artifacts such as `\u00e1` when they appear inside strings and then normalizes strings to Unicode Normalization Form C (NFC). NFC makes visually identical strings more stable for comparison while preserving their textual content.

For Reasoned-RAG and the reasoning-augmented trials of Mixed-SC, postprocessing removes the temporary reasoning/value interface. Each top-level object with `reasoning` and `value` fields is replaced by its `value`. The reasoning text is discarded before the final JSON object is returned and before any self-consistency aggregation is applied. If unwrapping fails because the expected structure is absent or malformed, the trial is invalid.

The prompt asks the model to use JSON `null` for unsupported information when null is allowed, but models may emit string literals such as `"null"`, `"NULL"`, or other casing variants instead. DRILLER converts these string markers to JSON `null` only at schema positions where `null` is valid. It does not globally replace all occurrences, because a non-nullable string field should not be silently changed into a null value.

In Mixed-SC, each trial independently performs prompt construction, strict generation, parsing, Unicode cleanup, optional unwrapping, and null coercion. Invalid trials are excluded from aggregation. The aggregator receives postprocessed JSON candidates in the original schema shape; after aggregation, nullable-string coercion may be applied again so selected values remain aligned with schema-valid null representation.

Thus, postprocessing in DRILLER is deterministic, schema-aware, and conservative. It makes generated objects cleaner and compatible with the expected interface while preserving the semantic content produced by the extractor.

8. Heterogeneous Self-Consistency

8.1. Motivation

Schema-guided extraction can vary across generations even when the output format is constrained. Candidate outputs may differ in whether they include a borderline entity, how they normalize a date, which enum label they choose, or whether they return JSON `null`. Some of this variance comes from sampling, but some comes from the interface shown to the model. For example, the reasoning-augmented prompt spends more output budget on field-local evidence assessment before committing to values.

In Mixed-SC we use this observation for heterogeneous self-consistency. Rather than sampling one prompt repeatedly, it combines Schema-RAG, which produces final values directly, with Reasoned-RAG, which temporarily produces `{reasoning, value}` objects and then unwraps them. The goal is to obtain candidates with potentially different error patterns. The final decision is not a flat vote over complete JSON strings; it is a recursive schema-aware aggregation over candidate objects.

8.2. Candidate Generation

Mixed-SC requests four extraction trials. Two follow the direct Schema-RAG style and two follow the reasoning-augmented Reasoned-RAG style. The intended ordering is interleaved. This keeps the candidate set balanced between the two extractor styles.

Each trial is a complete extraction run. It performs its own RAG retrieval, prompt construction, strict JSON generation, and deterministic postprocessing. Direct trials render examples and schemas in

the Schema-RAG format. Reasoning-augmented trials render examples and schemas in the temporary Reasoned-RAG format, so demonstrations contain reasoning/value fields during generation.

Postprocessing is applied before aggregation. Reasoning wrappers are removed, Unicode normalization and schema-aware null coercion are applied, and malformed outputs are rejected. The aggregation stage therefore receives final-shape JSON candidates in the original challenge schema, not raw model messages and not reasoning traces. Mixed-SC aggregates answers, not explanations.

Candidate validity is checked at the trial level. A candidate must parse as JSON and pass required unwrapping. Invalid candidates are excluded from aggregation. If all trials fail to produce valid postprocessed outputs, the system does not fabricate an extraction from partial information. The submitted plan requests four trials, with actual execution still subject to token and time budgets in the evaluation environment.

8.3. Schema-Aware Recursive Aggregation

Aggregation is recursive over the target schema. Complete JSON objects are not treated as indivisible strings. Instead, DRILLER descends through the schema and combines candidate values according to the expected type at each position. This matters because agreement has different meanings for closed labels, free text, nullable fields, nested objects, and arrays.

For closed scalar types, aggregation is close to exact voting. Enums, booleans, integers, and numeric values are grouped by canonical equality. The selected value comes from the group with strongest support. The aggregator does not synthesize new enum labels, invert booleans, average numbers, or create intermediate values.

Free-text strings are clustered rather than compared only by byte equality. Candidates may express the same grounded value with minor spelling, accent, punctuation, or phrasing differences. The default string path normalizes text and considers lexical overlap and character n-gram overlap. When a cluster is selected, the representative is a generated candidate value, typically a medoid, not a newly composed string.

Null values are treated as a separate candidate type. For nullable fields, JSON `null` is expected when the source text does not support a value. However, null does not win by default in ties: null wins only when its support is strictly greater than the best supported non-null value. This avoids turning uncertainty into absence while still preserving shared abstention when multiple trials independently return null.

Objects are aggregated field by field. For a root object, this allows one candidate to contribute stronger support for one field while another contributes a better-supported value elsewhere. The final object can therefore combine agreement patterns across the candidate set rather than selecting one whole candidate unchanged.

8.4. Array Aggregation

Arrays require separate handling because disagreement can involve order, omissions, duplicates, merged items, split items, or outliers. A majority vote over complete arrays would be brittle: two candidates may contain the same substantive items in different orders or differ by one peripheral object. DRILLER therefore aggregates arrays by clustering items across candidate arrays.

The array procedure first collects items from candidate arrays and compares them using schema-aware similarity. For arrays of simple values, item similarity follows the value type. For arrays of objects, object similarity compares corresponding fields with weights influenced by schema type. This allows two object items to align when they refer to the same entity, reaction, ingredient, event, or evidence span even if some attributes differ.

Near-duplicates from the same trial are handled cautiously. A single generated array may repeat one item or produce very similar variants. Such repetitions should not count as independent votes. The clustering process tracks trial indexes and prevents a non-exact cluster from accepting multiple values from the same trial, so one verbose generation cannot dominate support.

Arrays of objects may use identity-like fields to improve alignment. Many object-array schemas contain a field that behaves like an identifier: a name, title, mention, text span, ingredient, symptom, reaction, or source fragment. When such a field is detected with sufficient confidence, it guides item matching across trials. Two objects with the same reaction name but different impact labels should be compared as competing descriptions of one item, not as unrelated items. If no useful identity field is detected, the aggregator falls back to broader object similarity.

This alignment is important because object arrays often combine extraction and classification. The identity-like field tells the system which candidate objects correspond; remaining fields can then be compared or aggregated inside that item cluster. Imperfect identity detection is also a risk: a generic attribute may align unrelated objects, while failure to detect an identity field may split corresponding objects into duplicate clusters.

8.5. Recall-Biased Selection

After clustering array items, the aggregator selects a final array. The submitted design builds candidate final arrays from item clusters, including threshold-based arrays and a greedy MBR-style candidate. It compares possible arrays by their agreement with observed candidate arrays, using schema-aware item similarity rather than exact positional equality.

Selection is recall-biased when candidate arrays have similar scores. Among arrays within a tolerance of the best score, the selector prefers the longer array, then the higher-scoring one, then the earlier candidate. This favors retaining supported items when evidence from trials is close. The longer array is not accepted unconditionally; it must remain competitive under the consensus objective.

The design trades omission against over-extraction. A stricter selector would discard more borderline items, improving precision in some cases but risking missed entities or list elements. DRILLER’s recall-biased selector keeps items with enough support to remain near the best consensus score, which is useful for unordered lists, object arrays, and evidence collections where candidates often agree on a core set but vary on peripheral supported items.

8.6. Conservativeness and Failure Modes

Aggregation is conservative in that it selects among generated candidates rather than inventing semantic content. Exact scalar types are selected by canonical equality groups; non-exact string clusters choose a generated representative; arrays are built from generated item clusters. The aggregator does not use external knowledge, synthesize new enum labels, or create facts absent from all candidate outputs.

Self-consistency is still not a guarantee of correctness. If multiple trials share the same misunderstanding of a field description, attend to the same distractor mention, or overtrust an unsupported fact, aggregation can preserve that mistake with high support. Heterogeneous prompting reduces reliance on one prompt representation, but all trials still share the source text, target schema, local FSP corpus, model endpoint, and broad extraction rules.

Array aggregation has additional failure modes. The recall-biased selector may retain extra borderline items, and identity-field detection can misalign object items or fail to merge true matches. These risks are the cost of trying to preserve recall in schemas where missing list members are common and where item order is not meaningful.

Mixed-SC also increases inference cost. It requests four model calls rather than one, and two calls use reasoning-augmented generation, which can produce longer intermediate outputs before unwrapping. Its technical contribution is therefore not efficiency, but the combination of heterogeneous candidate generation with recursive schema-aware aggregation over final-shape JSON candidates.

9. Experimental Setup

Our submitted solution consists of three final pipelines: Schema-RAG, Reasoned-RAG, and Mixed-SC. These correspond to the submitted names enriched-schema-rag,

enriched-inline-reasoning-rag, and mixed-extractors-self-consistency-rag.

The final evaluation follows the official GenSIE shared-task protocol. GenSIE evaluates Spanish extraction from a source text, an instruction, and a target JSON Schema, with systems returning grounded schema-conformant JSON objects. Dataset construction, official scoring, ranking, model selection, and any efficiency reporting are defined by the GenSIE overview paper, which is treated here as the authoritative source for evaluation metadata [1].

DRILLER is compatible with the hosted evaluation through an OpenAI-style chat interface. The submitted pipelines do not hard-code a private model identity: the evaluator supplies the model name at runtime, and the system forwards that model name to the hosted inference endpoint. Exact official model identifiers, decoding settings, evaluation split details, and token accounting should be taken from organizer-provided metadata or final run logs.

10. Results and Analysis

At the time of writing, official shared-task scores, ranks, per-model results, and final token-use figures have not been released in citable form. They will be reported once available. We do not use exploratory local artifacts as final evidence here because they are partial, based on local splits, or associated with non-final variants; mixing them with official results would make the comparison unclear.

Until official outputs are available, analysis is limited to qualitative error categories that should be quantified later by inspecting aligned gold and system outputs with the source text and schema. These categories are not claims about observed official frequencies.

Unsupported inference is a central risk. A model may return a value that is plausible or true in the world but absent from the current text. This is especially relevant for nullable fields and for fields resembling common factual questions. Reasoned-RAG makes evidence assessment explicit during generation, and Schema-RAG emphasizes grounded extraction rules, but final analysis should measure whether unsupported values remain a source of false positives.

Array errors form another likely category. Systems may miss list items, duplicate objects, extract only salient examples, or include distractor mentions. Mixed-SC can retain items generated by at least one valid trial through clustering and recall-biased selection, but aggregation cannot recover an item absent from every candidate and may preserve extra borderline items.

Nullability errors should be separated from other value errors. They include returning `null` when evidence exists, returning a value when absence is correct, and emitting the string `"null"` instead of JSON `null`. DRILLER postprocesses only the representational case when the schema permits null; the first two are semantic grounding errors.

Variable schemas can also induce confusion between similar fields, such as creator versus director, release year versus setting year, ingredient versus allergen, or legal title versus affected article. Schema-readable prompting and field-aware retrieval are intended to reduce this risk, but final error analysis should check whether close field descriptions or distractor mentions still lead to swapped values.

Finally, aggregation introduces its own trade-offs. Mixed-SC can preserve shared mistakes when several trials converge on the same wrong interpretation, and its recall-biased array selector may favor inclusion when candidate arrays are close in score. Official analysis should therefore compare the three submitted pipelines by metric, cost, and error category without attributing component-level improvements unless supported by controlled ablations.

11. Limitations

DRILLER is designed for schema-variable extraction, and several limitations follow from that design. The first is dependence on a local few-shot prompting corpus. The demonstrations are synthetic and curated, which gives controlled coverage of nulls, populated values, enum choices, arrays, nested structures, and distractors, but also requires careful construction and review. Each case must keep the

instruction, source text, schema, output, and rationale mutually consistent. Errors or narrow coverage in the corpus can propagate into prompts.

Coverage remains incomplete because GenSIE schemas vary by instance. A demonstration family can represent one schema well while failing to cover later fields with different names, descriptions, nesting, or value regimes. Complementary examples reduce reliance on one output pattern, but they cannot guarantee that every private evaluation schema has a close analogue. If retrieved examples mostly show populated arrays, frequent enum labels, or non-null values, they may still bias the model toward overproduction.

Retrieval depends on schema information quality. Same-schema retrieval is useful when structurally matching examples exist, but it has limited value when the local corpus lacks the target schema. Field-level fallback is more flexible, yet it relies on names, descriptions, type representations, and coarse type classes. Sparse or generic descriptions can make unrelated fields appear similar, while equivalent fields with different wording can be missed. Type compatibility prevents many implausible alignments, but it can also exclude examples whose meanings are related despite different schema forms.

The reasoning/value wrapper in Reasoned-RAG is a cost trade-off. It can encourage explicit evidence assessment and clearer decisions between a value and `null`, but it increases output length, consumes context, and may raise latency or token cost. Because the submitted wrapper is top-level only, it does not give every nested field its own explicit evidence scaffold. This keeps the temporary schema manageable, but complex nested objects may still require decisions that are only indirectly guided by top-level reasoning.

Mixed-SC further increases inference cost by requesting four trials, including two reasoning-augmented calls. Runtime budgets may limit completed trials, and practical efficiency depends on the model, backend, decoding configuration, and official token accounting. Its value must therefore be judged jointly by extraction quality and cost.

Aggregation is heuristic. It votes over closed scalar values, clusters strings, aggregates objects field by field, clusters array items, and selects arrays with a recall-biased strategy. These rules are schema-aware and conservative in the sense that selected representatives come from generated candidates, but they are not optimality guarantees. Shared mistakes can be reinforced, borderline array items can be retained, and imperfect identity-field detection can misalign object arrays.

Finally, component-level effects remain empirical questions. The paper describes the submitted architecture, but claims about the separate impact of schema rendering, retrieval, same-schema prompting, reasoning wrappers, postprocessing, or heterogeneous self-consistency require controlled ablations. Results may also depend strongly on the hosted model’s structured-output behavior, context limits, multilingual extraction ability, and instruction following.

12. Conclusion

DRILLER submitted a modular schema-guided extraction system for GenSIE 2026. The system treats each instance as Spanish information extraction under a variable JSON Schema, where the central challenge is interpreting field semantics from the current instruction, schema, and source text while returning grounded schema-conformant JSON.

The three submitted pipelines are Schema-RAG, Reasoned-RAG, and Mixed-SC. All use retrieval-augmented prompting. The common extraction spine combines schema-readable Pydantic-style prompting, strict structured JSON generation, field-aware few-shot retrieval, and conservative postprocessing. Reasoned-RAG adds temporary top-level reasoning/value wrappers that are removed before final output. Mixed-SC combines direct and reasoning-augmented candidates through heterogeneous self-consistency and recursive schema-aware aggregation.

The main design principle is to treat the schema as a runtime semantic specification rather than only an output validator. Readable prompt schemas help the model interpret fields; strict generation constrains the response interface; retrieved examples illustrate field behavior; postprocessing normalizes representation-level artifacts without semantic repair; and aggregation reconciles candidate values

according to schema type.

Future work should strengthen both retrieval and aggregation. Learned retrieval could complement schema fingerprints and field-similarity heuristics, especially when field descriptions are sparse. Automatic validation of FSP cases would help detect inconsistent demonstrations before they enter prompts. Aggregation could benefit from better uncertainty estimates, learned item alignment, and schema-specific constraints. Controlled ablations are needed to measure component contributions, and robust array/object extraction remains a priority because omissions, duplicates, and field-alignment errors interact most strongly in nested structures.

Declaration on Generative AI

During the preparation of this manuscript, generative AI tools were used to assist with repository inspection, organization of technical notes, drafting, and language editing. The authors reviewed, edited, and approved the final content and take full responsibility for the submitted paper. The DRILLER system also uses synthetic and curated few-shot resources as part of its retrieval component; any generative assistance involved in preparing those resources was followed by manual review and validation.

References

- [1] GenSIE Organizers, Overview of GenSIE 2026: Spanish schema-guided information extraction, in: Proceedings of IberLEF 2026, 2026. To appear.
- [2] IberLEF Organizers, Overview of IberLEF 2026, in: Proceedings of IberLEF 2026, 2026. To appear.